

#1 Inews 遍历

Inews 遍历

bert-base

env-hardware:

env-software:

model:

acc = 0.84534782984012783

log detail:

roberta-wwm

env-hardware:

env-software:

model:

acc = 0.8418418418418419

log detail:

#2 bert-base

#3 env-hardware:

- Quadro RTX 8000
- VRAM: 48GiB
- Platinum 8160 96核128G
- 4TB SSD

#3 env-software:

- pytorch
- cuda9
- cudnn7 1.0.1
- python 3.6

#3 model:

- chinese_wwm_pytorch(bert-base)

#3 acc = 0.84534782984012783

#3 log detail:

```
1 (d1) root@1b20aad9b5b6:/data/baselines/classifier_pytorch# sh
2 run_classifier_inews.sh
3 Start training text classification on: with bert-base-chinese
4 id:10165102211
5 data exists
6 Finish download dataset.
7 Start running...
8 n_gpus: 4
9 4
10 01/12/2020 12:21:53 - WARNING - root - Process rank: -1, device: cuda,
    n_gpu: 4,
11 distributed training: False, 16-bits training: False
12 01/12/2020 12:21:53 - INFO - transformers.configuration_utils - loading
    configuration
13 file https://s3.amazonaws.com/models.huggingface.co/bert/bert-base-
    chinese-config.json
14 from cache at
```

```
15 /home/ubuntu/.cache/torch/transformers/8a3b1cfe5da58286e12a0f5d7d182b8d6ec
a88c08e26c332e
16 e3817548cf7e60a.0c16faba8be66db3f02805c912e4cf94d3c9cffc1f12fa1a39906f9270
f76d33
17 01/12/2020 12:21:53 - INFO - transformers.configuration_utils - Model
config {
18 "attention_probs_dropout_prob": 0.1,
19 "directionality": "bidi",
20 "finetuning_task": "inews",
21 "hidden_act": "gelu",
22 "hidden_dropout_prob": 0.1,
23 "hidden_size": 768,
24 "initializer_range": 0.02,
25 "intermediate_size": 3072,
26 "layer_norm_eps": 1e-12,
27 "max_position_embeddings": 512,
28 "num_attention_heads": 12,
29 "num_hidden_layers": 12,
30 "num_labels": 2,
31 "output_attentions": false,
32 "output_hidden_states": false,
33 "output_past": true,
34 "pooler_fc_size": 768,
35 "pooler_num_attention_heads": 12,
36 "pooler_num_fc_layers": 3,
37 "pooler_size_per_head": 128,
38 "pooler_type": "first_token_transform",
39 "pruned_heads": {},
40 "torchscript": false,
41 "type_vocab_size": 2,
42 "use_bfloat16": false,
43 "vocab_size": 21128
44 }
45 01/12/2020 12:21:54 - INFO - transformers.tokenization_utils - loading
file
46 https://s3.amazonaws.com/models.huggingface.co/bert/bert-base-chinese-
vocab.txt from
47 cache at
48 /home/ubuntu/.cache/torch/transformers/8a0c070123c1f794c42a29c6904beb7c1b8
715741e235bee0
49 4aca2c7636fc83f.9b42061518a39ca00b8b52059fd2bede8daa613f8a8671500e518a8c29
de8c00
50 01/12/2020 12:21:55 - INFO - transformers.modeling_utils - loading weights
file
51 https://s3.amazonaws.com/models.huggingface.co/bert/bert-base-chinese-
pytorch_model.bin
52 from cache at
53 /home/ubuntu/.cache/torch/transformers/b1b5e295889f2d0979ede9a78ad9cb5dc6a
0e25ab7f9417b3
54 15f0a2c22f4683d.929717ca66a3ba9eb9ec2f85973c6398c54c38a4faa464636a491d7a70
5f7eb6
55 01/12/2020 12:21:58 - INFO - transformers.modeling_utils - weights of
56 BertForSequenceClassification not initialized from pretrained model:
57 ['classifier.weight', 'classifier.bias']
58 01/12/2020 12:21:58 - INFO - transformers.modeling_utils - weights from
pretrained
59 model not used in BertForSequenceClassification: ['cls.predictions.bias',
60 'cls.predictions.transform.dense.weight',
'cls.predictions.transform.dense.bias',
61 'cls.predictions.decoder.weight', 'cls.seq_relationship.weight',
62 'cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.weight',
63 'cls.predictions.transform.LayerNorm.bias']
64 01/12/2020 12:22:03 - INFO - root - Training/evaluation parameters
65 Namespace(adam_epsilon=1e-08, cache_dir='', config_name='',
```

```

66 data_dir='/home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/CL
UEdatasets/afq
67 mc/', device=device(type='cuda'), do_eval=True, do_lower_case=True,
do_predict=False,
68 do_train=True, eval_all_checkpoints=False, fp16=False,
fp16_opt_level='O1',
69 gradient_accumulation_steps=1, learning_rate=2e-05, local_rank=-1,
logging_steps=2146,
70 max_grad_norm=1.0, max_seq_length=128, max_steps=-1,
71 model_name_or_path='bert-base-chinese', model_type='bert', n_gpu=4,
no_cuda=False,
72 num_train_epochs=3.0,
73 output_dir='/home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/
inews_output/b
74 ert', output_mode='classification', overwrite_cache=False,
overwrite_output_dir=True,
75 per_gpu_eval_batch_size=16, per_gpu_train_batch_size=16,
predict_checkpoints=0,
76 save_steps=2146, seed=42, server_ip='', server_port='', task_name='inews',
77 tokenizer_name='', warmup_proportion=0.1, weight_decay=0.01)
78 01/12/2020 12:22:03 - INFO - root - Loading features from cached file
79 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/CLUEdatasets
/inews/cached_
80 train_bert-base-chinese_128_inews
81 01/12/2020 12:22:04 - INFO - root - ***** Running training *****
82 01/12/2020 12:22:04 - INFO - root - Num examples = 34334
83 01/12/2020 12:22:04 - INFO - root - Num Epochs = 3
84 01/12/2020 12:22:04 - INFO - root - Instantaneous batch size per GPU = 16
85 01/12/2020 12:22:04 - INFO - root - Total train batch size (w. parallel,
distributed
86 & accumulation) = 64
87 01/12/2020 12:22:04 - INFO - root - Gradient Accumulation steps = 1
88 01/12/2020 12:22:04 - INFO - root - Total optimization steps = 1611
89 /home/ubuntu/anaconda3/lib/python3.7/site-
packages/torch.nn.parallel/_functions.py:61:
90 UserWarning: was asked to gather along dimension 0, but all input tensors
were scalars;
91 will instead unsqueeze and return a vector.
92 warnings.warn('was asked to gather along dimension 0, but all '
93 [Training] 537/537 [=====] 266.1ms/step loss:
0.4132
94 [Training] 537/537 [=====] 232.4ms/step loss:
0.5030
95 [Training] 537/537 [=====] 232.3ms/step loss:
0.4571
96 01/12/2020 12:28:37 - INFO - root - global_step = 1611, average loss =
97 0.47748216146789374
98 01/12/2020 12:28:37 - INFO - root - Saving model checkpoint to
99 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
/bert
100 01/12/2020 12:28:37 - INFO - transformers.configuration_utils -
Configuration saved in
101 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
/bert/config.j
102 son
103 01/12/2020 12:28:38 - INFO - transformers.modeling_utils - Model weights
saved in
104 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
/bert/pytorch_
105 model.bin
106 01/12/2020 12:28:38 - INFO - transformers.configuration_utils - loading
configuration
107 file
108 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
/bert/config.j

```

```
109 son
110 01/12/2020 12:28:38 - INFO - transformers.configuration_utils - Model
111 config {
112   "attention_probs_dropout_prob": 0.1,
113   "directionality": "bidi",
114   "finetuning_task": "inews",
115   "hidden_act": "gelu",
116   "hidden_dropout_prob": 0.1,
117   "hidden_size": 768,
118   "initializer_range": 0.02,
119   "intermediate_size": 3072,
120   "layer_norm_eps": 1e-12,
121   "max_position_embeddings": 512,
122   "num_attention_heads": 12,
123   "num_hidden_layers": 12,
124   "num_labels": 2,
125   "output_attentions": false,
126   "output_hidden_states": false,
127   "output_past": true,
128   "pooler_fc_size": 768,
129   "pooler_num_attention_heads": 12,
130   "pooler_num_fc_layers": 3,
131   "pooler_size_per_head": 128,
132   "pooler_type": "first_token_transform",
133   "pruned_heads": {},
134   "torchscript": false,
135   "type_vocab_size": 2,
136   "use_bfloat16": false,
137   "vocab_size": 21128
138 }
139 01/12/2020 12:28:38 - INFO - transformers.modeling_utils - loading weights
140 file
141 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output/
142 bert/pytorch_model.bin
143 01/12/2020 12:28:41 - INFO - transformers.tokenization_utils - Model name
144 '/home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output/
145 bert' not
146 found in model shortcut name list (bert-base-uncased, bert-large-uncased,
147 bert-base-cased, bert-large-cased, bert-base-multilingual-uncased,
148 bert-base-multilingual-cased, bert-base-chinese, bert-base-german-cased,
149 bert-large-uncased-whole-word-masking, bert-large-cased-whole-word-
150 masking,
151 bert-large-uncased-whole-word-masking-finetuned-squad,
152 bert-large-cased-whole-word-masking-finetuned-squad, bert-base-cased-
153 finetuned-mrpc,
154 bert-base-german-dbmz-cased, bert-base-german-dbmz-uncased). Assuming
155 '/home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output/
156 bert' is a
157 path or url to a directory containing tokenizer files.
158 01/12/2020 12:28:41 - INFO - transformers.tokenization_utils - loading
159 file
160 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output/
161 bert/vocab.txt
162 01/12/2020 12:28:41 - INFO - transformers.tokenization_utils - loading
163 file
164 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output/
165 bert/added_tokens.json
166 01/12/2020 12:28:41 - INFO - transformers.tokenization_utils - loading
167 file
168 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output/
169 bert/special_
170 tokens_map.json
```

```
161 01/12/2020 12:28:41 - INFO - transformers.tokenization_utils - loading
162 file
163 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
164 /bert/tokenize
165 r_config.json
166 01/12/2020 12:28:41 - INFO - transformers.tokenization_utils - Model name
167 '/home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_outpu
168 t/bert' not
169 found in model shortcut name list (bert-base-uncased, bert-large-uncased,
170 bert-base-cased, bert-large-cased, bert-base-multilingual-uncased,
171 bert-base-multilingual-cased, bert-base-chinese, bert-base-german-cased,
172 bert-large-uncased-whole-word-masking, bert-large-cased-whole-word-
173 masking,
174 bert-large-uncased-whole-word-masking-finetuned-squad,
175 bert-large-cased-whole-word-masking-finetuned-squad, bert-base-cased-
176 finetuned-mrpc,
177 bert-base-german-dbmz-cased, bert-base-german-dbmz-uncased). Assuming
178 '/home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_outpu
179 t/bert' is a
180 path or url to a directory containing tokenizer files.
181 01/12/2020 12:28:41 - INFO - transformers.tokenization_utils - loading
182 file
183 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
184 /bert/vocab.tx
185 t
186 01/12/2020 12:28:41 - INFO - transformers.tokenization_utils - loading
187 file
188 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
189 /bert/added_to
190 kens.json
191 01/12/2020 12:28:41 - INFO - transformers.tokenization_utils - loading
192 file
193 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
194 /bert/special_
195 tokens_map.json
196 01/12/2020 12:28:41 - INFO - transformers.tokenization_utils - loading
197 file
198 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
199 /bert/tokenize
200 r_config.json
201 01/12/2020 12:28:41 - INFO - root - Evaluate the following checkpoints:
202 ['/home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_outp
203 ut/bert']
204 01/12/2020 12:28:41 - INFO - transformers.configuration_utils - loading
205 configuration
206 file
207 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
208 /bert/config.j
209 son
210 01/12/2020 12:28:41 - INFO - transformers.configuration_utils - Model
211 config {
212 "attention_probs_dropout_prob": 0.1,
213 "directionality": "bidi",
214 "finetuning_task": "inews",
215 "hidden_act": "gelu",
216 "hidden_dropout_prob": 0.1,
217 "hidden_size": 768,
218 "initializer_range": 0.02,
219 "intermediate_size": 3072,
220 "layer_norm_eps": 1e-12,
221 "max_position_embeddings": 512,
222 "num_attention_heads": 12,
223 "num_hidden_layers": 12,
224 "num_labels": 2,
225 "output_attentions": false,
```

```

208 "output_hidden_states": false,
209 "output_past": true,
210 "pooler_fc_size": 768,
211 "pooler_num_attention_heads": 12,
212 "pooler_num_fc_layers": 3,
213 "pooler_size_per_head": 128,
214 "pooler_type": "first_token_transform",
215 "pruned_heads": {},
216 "torchscript": false,
217 "type_vocab_size": 2,
218 "use_bfloat16": false,
219 "vocab_size": 21128
220 }
221 01/12/2020 12:28:41 - INFO - transformers.modeling_utils - loading weights
file
222 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/inews_output
/bert/pytorch_
223 model.bin
224 01/12/2020 12:28:43 - INFO - root - Loading features from cached file
225 /home/ubuntu/CLUE/baselines/models_pytorch/classifier_pytorch/CLUEdatasets
/inews/cached_
226 dev_bert-base-chinese_128_inews
227 01/12/2020 12:28:43 - INFO - root - ***** Running evaluation *****
228 [Evaluating] 68/68 [=====] 89.1ms/step
229 01/12/2020 12:28:49 - INFO - root - Num examples = 4316
230 01/12/2020 12:28:49 - INFO - root - Batch size = 64
231 01/12/2020 12:28:49 - INFO - root - ***** Eval results *****
232 01/12/2020 12:28:49 - INFO - root - dev: acc = 0.84534782984012783

```

#2 roberta-wwm

#3 env-hardware:

- Quadro RTX 8000
- VRAM: 48GiB
- Platinum 8160 96核128G
- 4TB SSD

#3 env-software:

- pytorch
- cuda9
- cudnn7 1.0.1
- python 3.6

#3 model:

- chinese_roberta_wwm_ext_pytorch

#3 acc = 0.8418418418418419

#3 log detail:

```

1 (d1) root@1b20aad9b5b6:/data/baselines/classifier_pytorch# sh
run_classifier_inews.sh
2 /usr/local/miniconda3/envs/d1/lib/python3.6/site-
packages/sklearn/externals/joblib/externals/cloudpickle/cloudpickle.py:47:
DeprecationWarning: the imp module is deprecated in favour of importlib;
see the module's documentation for alternative uses

```

```

3     import imp
4 01/14/2020 14:53:47 - WARNING - root - Process rank: -1, device: cuda,
n_gpu: 1, distributed training: False, 16-bits training: False
5 01/14/2020 14:53:47 - INFO - transformers.configuration_utils - loading
configuration file
/data/baselines/classifier_pytorch/prev_trained_model/chinese_roberta_wwm_
ext_pytorch/config.json
6 01/14/2020 14:53:47 - INFO - transformers.configuration_utils - Model
config {
7     "attention_probs_dropout_prob": 0.1,
8     "directionality": "bidi",
9     "finetuning_task": "inews",
10    "hidden_act": "gelu",
11    "hidden_dropout_prob": 0.1,
12    "hidden_size": 768,
13    "initializer_range": 0.02,
14    "intermediate_size": 3072,
15    "layer_norm_eps": 1e-12,
16    "max_position_embeddings": 512,
17    "num_attention_heads": 12,
18    "num_hidden_layers": 12,
19    "num_labels": 3,
20    "output_attentions": false,
21    "output_hidden_states": false,
22    "output_past": true,
23    "pooler_fc_size": 768,
24    "pooler_num_attention_heads": 12,
25    "pooler_num_fc_layers": 3,
26    "pooler_size_per_head": 128,
27    "pooler_type": "first_token_transform",
28    "pruned_heads": {},
29    "torchscript": false,
30    "type_vocab_size": 2,
31    "use_bfloat16": false,
32    "vocab_size": 21128
33 }
34
35 01/14/2020 14:53:47 - INFO - transformers.tokenization_utils - Model
name
'/data/baselines/classifier_pytorch/prev_trained_model/chinese_roberta_wwm_
_ext_pytorch' not found in model shortcut name list(bert-base-uncased,
bert-large-uncased, bert-base-cased, bert-large-cased, bert-base-
multilingual-uncased, bert-base-multilingual-cased, bert-base-chinese,
bert-base-german-cased, bert-large-uncased-whole-word-masking, bert-large-
cased-whole-word-masking, bert-large-uncased-whole-word-masking-finetuned-
squad, bert-large-cased-whole-word-masking-finetuned-squad, bert-base-
cased-finetuned-mrpc, bert-base-german-dbmdz-cased, bert-base-german-
dbmdz-uncased). Assuming
'/data/baselines/classifier_pytorch/prev_trained_model/chinese_roberta_wwm_
_ext_pytorch' is a path or url to a directory containing tokenizer files.
36 01/14/2020 14:53:47 - INFO - transformers.tokenization_utils - Didn't
find file
/data/baselines/classifier_pytorch/prev_trained_model/chinese_roberta_wwm_
_ext_pytorch/added_tokens.json. We won't load it.
37 01/14/2020 14:53:48 - INFO - transformers.tokenization_utils - Didn't
find file
/data/baselines/classifier_pytorch/prev_trained_model/chinese_roberta_wwm_
_ext_pytorch/special_tokens_map.json. We won't load it.
38 01/14/2020 14:53:48 - INFO - transformers.tokenization_utils - Didn't
find file
/data/baselines/classifier_pytorch/prev_trained_model/chinese_roberta_wwm_
_ext_pytorch/tokenizer_config.json. We won't load it.

```

```

39 01/14/2020 14:53:48 - INFO - transformers.tokenization_utils - loading
    file
    /data/baselines/classifier_pytorch/prev_trained_model/chinese_roberta_wwm_
    ext_pytorch/vocab.txt
40 01/14/2020 14:53:48 - INFO - transformers.tokenization_utils - loading
    file None
41 01/14/2020 14:53:48 - INFO - transformers.tokenization_utils - loading
    file None
42 01/14/2020 14:53:48 - INFO - transformers.tokenization_utils - loading
    file None
43 01/14/2020 14:53:48 - INFO - transformers.modeling_utils - loading
    weights file
    /data/baselines/classifier_pytorch/prev_trained_model/chinese_roberta_wwm_
    ext_pytorch/pytorch_model.bin
44 01/14/2020 14:54:12 - INFO - transformers.modeling_utils - weights of
    BertForSequenceClassification not initialized from pretrained model:
    ['classifier.weight', 'classifier.bias']
45 01/14/2020 14:54:12 - INFO - transformers.modeling_utils - weights from
    pretrained model not used in BertForSequenceClassification:
    ['cls.predictions.bias', 'cls.predictions.transform.dense.weight',
    'cls.predictions.transform.dense.bias',
    'cls.predictions.transform.LayerNorm.weight',
    'cls.predictions.transform.LayerNorm.bias',
    'cls.predictions.decoder.weight', 'cls.seq_relationship.weight',
    'cls.seq_relationship.bias']
46 01/14/2020 14:54:14 - INFO - root - Training/evaluation parameters
    Namespace(adam_epsilon=1e-08, cache_dir='', config_name='',
    data_dir='/data/baselines/classifier_pytorch/chineseGLUEdatasets/inews/',
    device=device(type='cuda'), do_eval=True, do_lower_case=True,
    do_predict=False, do_train=True, eval_all_checkpoints=False, fp16=False,
    fp16_opt_level='O1', gradient_accumulation_steps=1, learning_rate=2e-05,
    local_rank=-1, logging_steps=670, max_grad_norm=1.0, max_seq_length=512,
    max_steps=-1,
    model_name_or_path='/data/baselines/classifier_pytorch/prev_trained_model/
    chinese_roberta_wwm_ext_pytorch', model_type='roberta', n_gpu=1,
    no_cuda=False, num_train_epochs=4.0,
    output_dir='/data/baselines/classifier_pytorch/outputs/inews_output/robert
    a', output_mode='classification', overwrite_cache=False,
    overwrite_output_dir=True, per_gpu_eval_batch_size=8,
    per_gpu_train_batch_size=8, save_steps=670, seed=42, server_ip='',
    server_port='', task_name='inews', tokenizer_name='',
    warmup_proportion=0.1, weight_decay=0.01)
47 01/14/2020 14:54:15 - INFO - root - Loading features from cached file
    /data/baselines/classifier_pytorch/chineseGLUEdatasets/inews/cached_train_
    chinese_roberta_wwm_ext_pytorch_512_inews
48 01/14/2020 14:54:18 - INFO - root - ***** Running training *****
49 01/14/2020 14:54:18 - INFO - root - Num examples = 5355
50 01/14/2020 14:54:18 - INFO - root - Num Epochs = 4
51 01/14/2020 14:54:18 - INFO - root - Instantaneous batch size per GPU =
    8
52 01/14/2020 14:54:18 - INFO - root - Total train batch size (w.
    parallel, distributed & accumulation) = 8
53 01/14/2020 14:54:18 - INFO - root - Gradient Accumulation steps = 1
54 01/14/2020 14:54:18 - INFO - root - Total optimization steps = 2680
55 [Training] 669/670 [=====>.] - ETA: 0s loss:
    0.6268 01/14/2020 14:58:09 - INFO - root - Creating features from
    dataset file at
    /data/baselines/classifier_pytorch/chineseGLUEdatasets/inews/
56 01/14/2020 14:58:10 - INFO - processors.glue - writing example 0
57 01/14/2020 14:58:10 - INFO - processors.glue - *** Example ***
58 01/14/2020 14:58:10 - INFO - processors.glue - guid: dev-1

```



```

59 01/14/2020 14:58:10 - INFO - processors.glue - input_ids: 101 2526 2360
872 1962 8024 6435 7309 2769 1762 4397 4263 5381 677 4692 1168 671 702
2552 811 4638 1996 2023 8024 1377 3221 5381 4991 119 119119 142 2526 2360
8728 102 4761 6399 772 3326 3221 1920 2157 6963 7478 2382 1068 3800 4638
6413 75798024 3684 4994 2769 812 6716 1905 754 671 702 6158 2255 2181 1259
1741 4638 4852 833 8024 6929 720 1762 4761 6399 772 3326 6401 6390 704
8024 6963 3300 1525 763 1372 966 2533 2769 812 3800 2692 1450 8024 4761
6399 772 3326 6401 6390 3198 3126 3221 1914 719 8024 8024 4761 6399 772
3326 4509 6435 1453 3309 3221 1914 719 8024 809 1350 4761 6399 772 3326
6401 6390 2418 6401 3332 3160 3300 1525 763 1450 102 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 00 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
60 01/14/2020 14:58:10 - INFO - processors.glue - attention_mask: 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 11 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 00 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
61 01/14/2020 14:58:10 - INFO - processors.glue - token_type_ids: 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 11 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
62 01/14/2020 14:58:10 - INFO - processors.glue - label: 1 (id = 1)
63 01/14/2020 14:58:10 - INFO - processors.glue - input length: 136
64 01/14/2020 14:58:10 - INFO - processors.glue - *** Example ***
65 01/14/2020 14:58:10 - INFO - processors.glue - guid: dev-2

```


[illegible]

[illegible]

[illegible]

```
97 01/14/2020 14:58:38 - INFO - root - ***** Eval results *****
98 01/14/2020 14:58:38 - INFO - root - acc = 0.8388388388388388
99 01/14/2020 14:58:40 - INFO - transformers.configuration_utils -
Configuration saved in
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-670/config.json
100 01/14/2020 15:00:23 - INFO - transformers.modeling_utils - Model weights
saved in
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-670/pytorch_model.bin
101 01/14/2020 15:00:23 - INFO - root - Saving model checkpoint to
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-670
102 [Training] 670/670 [=====] 546.6ms/step loss:
0.0379
103 [Training] 669/670 [=====>.] - ETA: 0s loss:
0.1278 01/14/2020 15:04:14 - INFO - root - Loading features from cached
file
/data/baselines/classifier_pytorch/chineseGLUEdatasets/inews/cached_dev_ch
inese_roberta_wwm_ext_pytorch_512_inews
104 01/14/2020 15:04:15 - INFO - root - ***** Running evaluation *****
105 01/14/2020 15:04:15 - INFO - root - Num examples = 999
106 01/14/2020 15:04:15 - INFO - root - Batch size = 8
107 [Evaluating] 125/125 [=====] 115.5ms/step
108 01/14/2020 15:04:29 - INFO - root - ***** Eval results *****
109 01/14/2020 15:04:29 - INFO - root - acc = 0.8278278278278278
110 01/14/2020 15:04:31 - INFO - transformers.configuration_utils -
Configuration saved in
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-1340/config.json
111 01/14/2020 15:06:14 - INFO - transformers.modeling_utils - Model weights
saved in
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-1340/pytorch_model.bin
112 01/14/2020 15:06:15 - INFO - root - Saving model checkpoint to
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-1340
113 [Training] 670/670 [=====] 524.6ms/step loss:
0.2891
114 [Training] 669/670 [=====>.] - ETA: 0s loss:
0.0177 01/14/2020 15:10:06 - INFO - root - Loading features from cached
file
/data/baselines/classifier_pytorch/chineseGLUEdatasets/inews/cached_dev_ch
inese_roberta_wwm_ext_pytorch_512_inews
115 01/14/2020 15:10:07 - INFO - root - ***** Running evaluation *****
116 01/14/2020 15:10:07 - INFO - root - Num examples = 999
117 01/14/2020 15:10:07 - INFO - root - Batch size = 8
118 [Evaluating] 125/125 [=====] 115.6ms/step
119 01/14/2020 15:10:21 - INFO - root - ***** Eval results *****
120 01/14/2020 15:10:21 - INFO - root - acc = 0.8458458458458459
121 01/14/2020 15:10:23 - INFO - transformers.configuration_utils -
Configuration saved in
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-2010/config.json
122 01/14/2020 15:12:04 - INFO - transformers.modeling_utils - Model weights
saved in
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-2010/pytorch_model.bin
123 01/14/2020 15:12:05 - INFO - root - Saving model checkpoint to
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-2010
124 [Training] 670/670 [=====] 522.2ms/step loss:
0.0023
```

```

125 [Training] 669/670 [=====>.] - ETA: 0s loss:
0.4203 01/14/2020 15:15:55 - INFO - root - Loading features from cached
file
/data/baselines/classifier_pytorch/chineseGLUEdatasets/inews/cached_dev_ch
inese_roberta_wwm_ext_pytorch_512_inews
126 01/14/2020 15:15:57 - INFO - root - ***** Running evaluation *****
127 01/14/2020 15:15:57 - INFO - root - Num examples = 999
128 01/14/2020 15:15:57 - INFO - root - Batch size = 8
129 [Evaluating] 125/125 [=====] 115.7ms/step
130 01/14/2020 15:16:12 - INFO - root - ***** Eval results *****
131 01/14/2020 15:16:12 - INFO - root - acc = 0.8418418418418419
132 01/14/2020 15:16:13 - INFO - transformers.configuration_utils -
Configuration saved in
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-2680/config.json
133 01/14/2020 15:17:55 - INFO - transformers.modeling_utils - Model weights
saved in
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-2680/pytorch_model.bin
134 01/14/2020 15:17:56 - INFO - root - Saving model checkpoint to
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/checkpoint
-2680
135 [Training] 670/670 [=====] 524.4ms/step loss:
0.0022
136 01/14/2020 15:17:57 - INFO - root - global_step = 2680, average loss =
0.31527684231179265
137 01/14/2020 15:17:57 - INFO - root - Saving model checkpoint to
/data/baselines/classifier_pytorch/outputs/inews_output/roberta
138 01/14/2020 15:17:57 - INFO - transformers.configuration_utils -
Configuration saved in
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/config.js
on
139 01/14/2020 15:19:41 - INFO - transformers.modeling_utils - Model weights
saved in
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/pytorch_m
odel.bin
140 01/14/2020 15:19:44 - INFO - transformers.configuration_utils - loading
configuration file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/config.js
on
141 01/14/2020 15:19:45 - INFO - transformers.configuration_utils - Model
config {
142 "attention_probs_dropout_prob": 0.1,
143 "directionality": "bidi",
144 "finetuning_task": "inews",
145 "hidden_act": "gelu",
146 "hidden_dropout_prob": 0.1,
147 "hidden_size": 768,
148 "initializer_range": 0.02,
149 "intermediate_size": 3072,
150 "layer_norm_eps": 1e-12,
151 "max_position_embeddings": 512,
152 "num_attention_heads": 12,
153 "num_hidden_layers": 12,
154 "num_labels": 3,
155 "output_attentions": false,
156 "output_hidden_states": false,
157 "output_past": true,
158 "pooler_fc_size": 768,
159 "pooler_num_attention_heads": 12,
160 "pooler_num_fc_layers": 3,
161 "pooler_size_per_head": 128,
162 "pooler_type": "first_token_transform",
163 "pruned_heads": {},
164 "torchscript": false,

```

```

165     "type_vocab_size": 2,
166     "use_bfloat16": false,
167     "vocab_size": 21128
168 }
169
170 01/14/2020 15:19:45 - INFO - transformers.modeling_utils - loading
weights file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/pytorch_model.bin
171 01/14/2020 15:20:13 - INFO - transformers.tokenization_utils - Model
name '/data/baselines/classifier_pytorch/outputs/inews_output/roberta' not
found in model shortcut name list (bert-base-uncased, bert-large-uncased,
bert-base-cased, bert-large-cased, bert-base-multilingual-uncased, bert-
base-multilingual-cased, bert-base-chinese, bert-base-german-cased, bert-
large-uncased-whole-word-masking, bert-large-cased-whole-word-masking,
bert-large-uncased-whole-word-masking-finetuned-squad, bert-large-cased-
whole-word-masking-finetuned-squad, bert-base-cased-finetuned-mrpc, bert-
base-german-dbmdz-cased, bert-base-german-dbmdz-uncased). Assuming
'/data/baselines/classifier_pytorch/outputs/inews_output/roberta' is a
path or url to a directory containing tokenizer files.
172 01/14/2020 15:20:14 - INFO - transformers.tokenization_utils - loading
file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/vocab.txt
173 01/14/2020 15:20:14 - INFO - transformers.tokenization_utils - loading
file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/added_tokens.json
174 01/14/2020 15:20:14 - INFO - transformers.tokenization_utils - loading
file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/special_tokens_map.json
175 01/14/2020 15:20:14 - INFO - transformers.tokenization_utils - loading
file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/tokenizer_config.json
176 01/14/2020 15:20:14 - INFO - transformers.tokenization_utils - Model
name '/data/baselines/classifier_pytorch/outputs/inews_output/roberta' not
found in model shortcut name list (bert-base-uncased, bert-large-uncased,
bert-base-cased, bert-large-cased, bert-base-multilingual-uncased, bert-
base-multilingual-cased, bert-base-chinese, bert-base-german-cased, bert-
large-uncased-whole-word-masking, bert-large-cased-whole-word-masking,
bert-large-uncased-whole-word-masking-finetuned-squad, bert-large-cased-
whole-word-masking-finetuned-squad, bert-base-cased-finetuned-mrpc, bert-
base-german-dbmdz-cased, bert-base-german-dbmdz-uncased). Assuming
'/data/baselines/classifier_pytorch/outputs/inews_output/roberta' is a
path or url to a directory containing tokenizer files.
177 01/14/2020 15:20:14 - INFO - transformers.tokenization_utils - loading
file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/vocab.txt
178 01/14/2020 15:20:14 - INFO - transformers.tokenization_utils - loading
file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/added_tokens.json
179 01/14/2020 15:20:14 - INFO - transformers.tokenization_utils - loading
file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/special_tokens_map.json
180 01/14/2020 15:20:14 - INFO - transformers.tokenization_utils - loading
file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/tokenizer_config.json
181 01/14/2020 15:20:15 - INFO - root - Evaluate the following checkpoints:
['/data/baselines/classifier_pytorch/outputs/inews_output/roberta']

```



```

182 01/14/2020 15:20:15 - INFO - transformers.configuration_utils - loading
configuration file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/config.js
on
183 01/14/2020 15:20:15 - INFO - transformers.configuration_utils - Model
config {
184     "attention_probs_dropout_prob": 0.1,
185     "directionality": "bidi",
186     "finetuning_task": "inews",
187     "hidden_act": "gelu",
188     "hidden_dropout_prob": 0.1,
189     "hidden_size": 768,
190     "initializer_range": 0.02,
191     "intermediate_size": 3072,
192     "layer_norm_eps": 1e-12,
193     "max_position_embeddings": 512,
194     "num_attention_heads": 12,
195     "num_hidden_layers": 12,
196     "num_labels": 3,
197     "output_attentions": false,
198     "output_hidden_states": false,
199     "output_past": true,
200     "pooler_fc_size": 768,
201     "pooler_num_attention_heads": 12,
202     "pooler_num_fc_layers": 3,
203     "pooler_size_per_head": 128,
204     "pooler_type": "first_token_transform",
205     "pruned_heads": {},
206     "torchscript": false,
207     "type_vocab_size": 2,
208     "use_bfloat16": false,
209     "vocab_size": 21128
210 }
211
212 01/14/2020 15:20:15 - INFO - transformers.modeling_utils - loading
weights file
/data/baselines/classifier_pytorch/outputs/inews_output/roberta/pytorch_mo
del.bin
213 01/14/2020 15:20:41 - INFO - root - Loading features from cached file
/data/baselines/classifier_pytorch/chineseGLUEdatasets/inews/cached_dev_ch
inese_roberta_wwm_ext_pytorch_512_inews
214 01/14/2020 15:20:41 - INFO - root - ***** Running evaluation *****
215 01/14/2020 15:20:41 - INFO - root - Num examples = 999
216 01/14/2020 15:20:41 - INFO - root - Batch size = 8
217 [Evaluating] 125/125 [=====] 112.5ms/step
218 01/14/2020 15:20:56 - INFO - root - ***** Eval results *****
219 01/14/2020 15:20:56 - INFO - root - acc = 0.8418418418418419

```

复现成功。